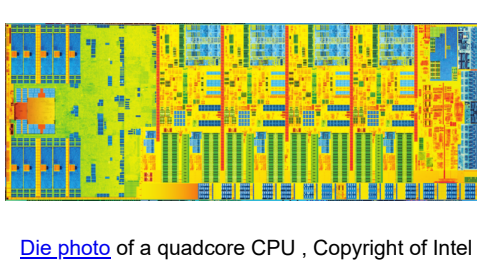


MICROARCHITECTURE CHEAT SHEET

X86 CPUs & Performance



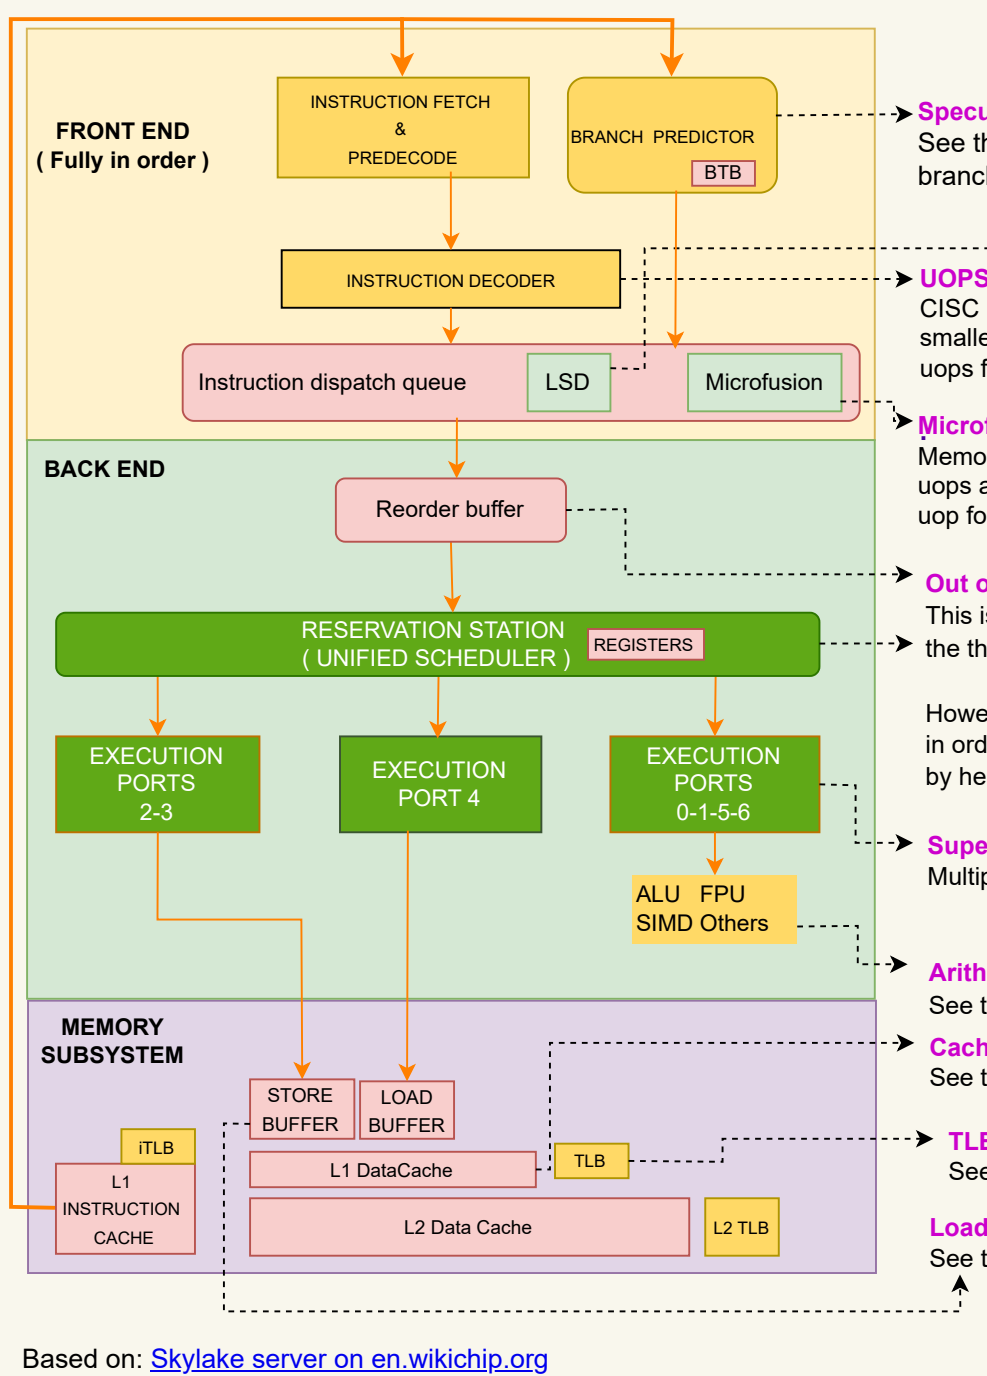
LAST UPDATE DATE : 11 NOV 2022

FOR LATEST VERSION: www.github.com/akhilmicroarchitecture/cheatsheet

AUTHOR: AKIN OCAL akin_ocal@hotmail.com

PIPELINE REALM : INSIDE AN INDIVIDUAL CORE

A SIMPLIFIED OVERVIEW (BASED ON INTEL SKYLAKE)



Based on: [Skylake server on en.wikichip.org](https://en.wikichip.org/wiki/Skylake_server)

Speculative execution
See the branch predictor realm for branch predictor, BTB and LSD.

UCOPS/MicroOps
CISC instructions are split into smaller RISC instructions called as UCOPS for more throughput.

Microfusion
Memory-write and read-modify ops are combined into a single uop for better throughput.

Out of order execution
This is done at the backend to improve the throughput.

However the results have to be output in order. So reordering at the end done by help of the reorder buffer.

Super scalar execution
Multiple uops executed in parallel.

Arithmetic operations
See the arithmetic realm.

Caches
See the cache memory realm.

TLBs
See the virtual memory realm.

Load/Store buffers
See the load-store realm.

PIPELINE PARALLELISM & PERFORMANCE

Pipeline diagrams : The diagrams below in the following topics are outputs from an online microarchitecture analysis tool [wikichip](https://en.wikichip.org/wiki/Tools) and they represent parallel execution through cycles.

Rows are multiple instructions being executed at the same time.

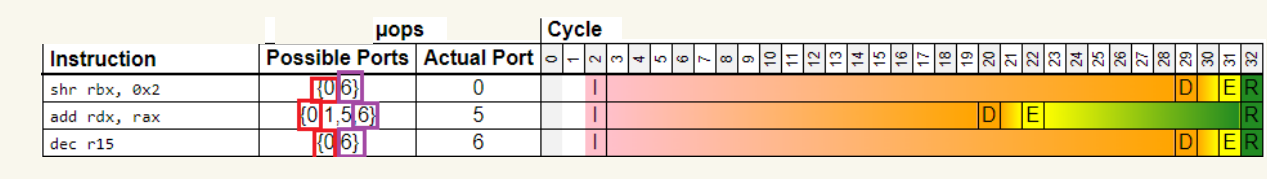
Cumulative display how instruction state changes through cycles.

IPC : As for pipeline performance, typically IPC is used. It stands for "Instructions per cycle". A higher IPC value usually means a better throughput.

You can measure IPC with perf : <https://perf.wiki.kernel.org/index.php/Tutorial>

Rate of retired instructions : Apart from IPC, number of retired instructions should be checked. Retired instructions are not committed/fetched as they were wrongly speculated. On the other hand executed instructions are the ones which were fetched. Therefore a high rate of retired instructions indicates low branch prediction rate.

CONTENTION FOR EXECUTION PORTS IN THE PIPELINE

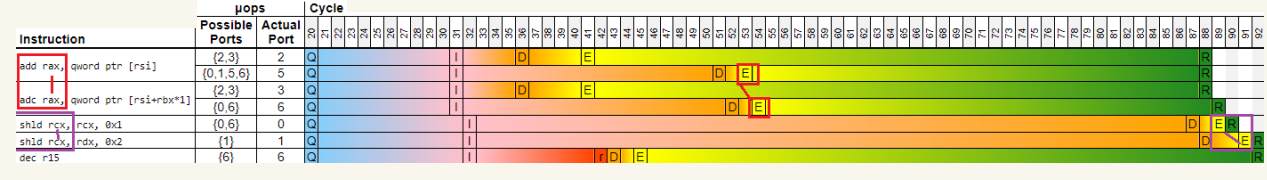


In the example above, all instructions are working on different registers, but SHR, ADD, DEC instructions are competing for ports 0 and 6. SHR and DEC are getting executed after ADD instruction.

Also notice that there is longer time between (Executed) and (Retired) states of instruction ADD as retirement has to be done in-order whereas execution is out-of-order.

Reference : [Dennis Bakhtov's article](https://en.wikichip.org/wiki/Tools)

INSTRUCTION STALLS DUE TO DATA DEPENDENCY



In the example above, there are 2 dependency chains, each marked with a different colour. In the first red coloured one, 2 instructions are competing for RAX register and notice that the second instruction gets executed after the first one. And the same applies to the 2nd purple pair.

Reference : [Dennis Bakhtov's article](https://en.wikichip.org/wiki/Tools)

RTDSCP INSTRUCTION FOR MEASUREMENTS

TSC (time stamp counter) is a special register that counts CPU cycles. RTDSCP can be used to read the TSC value which then can be used for measurements. It can also avoid out-of-order execution effects to a degree : [From Intel Software Developer's Manual Volume 4.1, April 2022](https://en.wikichip.org/wiki/Tools)

Intel's Way to benchmark code execution times : whitelap has details of using RTDSCP instruction.

AMD Programmers Manual Vol3 states : RTDSCP forces all instructions to retire before reading the time-stamp counter

ESTIMATING INSTRUCTION LATENCIES

You can use Agner Fog's <https://en.wikichip.org/wiki/Tools> to find out instructions' reciprocal throughputs (clock cycle per instruction). As an example, reciprocal throughput of instruction RTDSCP is 32 on Skylake microarchitecture :

> 1 cycle (4.5GHz) (highest frequency on Skylake) is 0.22 nanoseconds

> 32*0.22=7.04 nanoseconds

So its resolution estimation is about 7 nanoseconds on a 4.5 GHz Skylake CPU. You have to recalculate it for different microarchitectures and clock speeds.

HYPERTHREADING / SIMULTANEOUS MULTITHREADING

Hyperthreading name is used by Intel and it is called as "Simultaneous multithreading" by AMD. In both resources including CPUs and execution units are shared Agner Fog's <https://en.wikichip.org/wiki/Tools> has "multithreading" sections for each of Intel and AMD microarchitectures.

Regarding using it, if your app is data-intensive, hyperthreading won't help. Therefore it can be disabled via BIOS settings.

In general, it moves the control of resources from software to hardware and that is usually not desired for performance critical applications.

DYNAMIC FREQUENCIES

Modern CPUs employ dynamic frequency scaling which means there is a min and a max frequency per CPU core.

ACPI : ACPI defines multiple power states and modern CPUs implement those. P-State is for performance and C-states are for energy efficiency.

Intel has various **frequency options** and the most well known is TurboBoost. On AMD side there is **TurboCore**. You can use those to maximise the CPU usage.

Number of active cores & SIMD AVX2/SSE12 on Intel CPUs : Intel's power management policies are complex. See the arithmetic and the multicore realms as number of active cores and some of AVX2/SSE12 extensions also may affect the frequency while in TurboBoost.

LOAD STORE REALM

LOAD & STORE BUFFERS

Load and store buffers allow CPU to do out-of-order execution on loads and stores by decoupling speculative execution and committing the results to the cache memory.

Reference : https://en.wikichip.org/wiki/Memory_disambiguation

STORE-TO-LOAD FORWARDING

Using buffers for stores and loads to support out of order execution leads to a data synchronization issue. That issue is described in en.wikichip.org/wiki/Memory_disambiguation Store-to-load forwarding. As a solution, CPU can forward a memory store operation to a following load, if they are both operating on the same address.

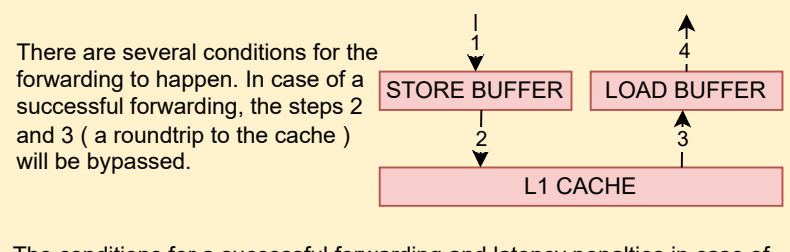
An example store and load sequence :

```
mov [eax], ecx ; STORE. Write the value of ECX register to the memory address which is stored in EAX register
mov ecx, [eax] ; LOAD. Read the value from that memory address
mov ecx, [eax] ; (which was just used) and write it to ECX register
```

STORE-TO-LOAD FORWARDING & LHS & PERFORMANCE

Based on [Intel Optimization Manual 3.6.4](https://en.wikichip.org/wiki/Memory_disambiguation), store-to-load forwarding may improve combined latency of those 2 operations. The reason is not specified however it is potentially LHS (Load-Hit-Store) problem in which the penalty is a round trip to the cache memory.

<https://en.wikichip.org/wiki/Load-Hit-Store>



There are several conditions for the forwarding to happen. In case of a successful forwarding, the steps 2 and 3 (a roundtrip to the cache) will be bypassed.

The conditions for a successful forwarding and latency penalties in case of non-forwarding can be found in Agner Fog's [microarchitecture book](https://en.wikichip.org/wiki/Memory_disambiguation).

What would happen without forwarding? : In the past, game consoles (PlayStation3) and Xbox360 had PowerPC based processors which used in-order execution rather than out-of-order execution. Therefore developers had to separately handle LHS by using `memory_order_seq_cst` keyword and other methods : [Dmitry Ryukin's article](https://en.wikichip.org/wiki/Memory_disambiguation)

ARITHMETIC REALM

ARITHMETIC INSTRUCTION LATENCIES

You can see a set of arithmetic operations from fast to slow below.

The clock cycles are based on Agner Fog's <https://en.wikichip.org/wiki/Tools> and Skylake architecture on 64 bit registers.

Bitwise operations : integer arithmetic : 0.25 to 1 clock cycle

Floating point multiplication : about 3 clock cycles

Floating point division : about 10 clock cycles

Integer division : 24-30 clock cycles

FLOATING POINTS

X86 uses [IEEE 754](https://en.wikichip.org/wiki/Tools) standard for floating points. A 32 bit floating point consists of 3 parts in the memory layout. Below you can see all bits of 1284 5879 FP number. Using <https://en.wikichip.org/wiki/Tools> as visualizer :

A floating point's value is calculated as : mantissa * 2^{exponent}

IEEE754 also defines **denormal numbers**. They are very small / near zero numbers.

As floating points are approximations, denormal numbers are needed to avoid an undefined case of `while (a-b) { ... }` without denormal the code to the right would invoke a divide-by-zero exception.

Reference : [Bruce Dawson's article](https://en.wikichip.org/wiki/Tools)

Based on Agner Fog's [microarchitecture book](https://en.wikichip.org/wiki/Tools), Intel CPUs have a penalty for denormal numbers, for ex: 129 clock cycles on Skylake. They also can be turned off on Intel CPUs.

As for AMD side, the recent Zen architecture CPUs seemingly don't have the same performance degradation.

X86 EXTENSIONS

X86 extensions are specialised instructions. They have various categories from [vectorization](https://en.wikichip.org/wiki/Tools) to [network operations](https://en.wikichip.org/wiki/Tools).

[Intel Intrinsics Guide](https://en.wikichip.org/wiki/Tools) is a good page to explore those extensions.

SSE (Streaming SIMD Extensions) is one of the most important ones. **SIMD** stands for "single instruction multiple data". SIMD instructions use wider registers to execute more work in a single go.

In the example above, an array of 4 integers (11 to 44) are added to another array of integers (11 to 44). The result is also an array of sums (11 to 44). In this example, 4 add operations are executed by a single instruction.

They play key role in complex vectorisation optimisations : [GCC auto vectorization](https://en.wikichip.org/wiki/Tools)

Apart from arithmetic operations, they can be utilised for string operations as well.

A SIMD based JSON parser : <https://en.wikichip.org/wiki/Tools>

X86 EXTENSIONS : SIMD DETAILS

The most recent SIMD instruction sets and their corresponding registers are :

AVX1 : 128 bits, XMM registers

AVX2 : 256 bits, YMM registers

AVX512 : 512 bits, ZMM registers

Note about AVX512 : AMD started to implement it with Zen4.

As for programming, there are also wider data types. The data type diagrams below are for 128 bit AVX.

For details : [Daniel Lemire's article](https://en.wikichip.org/wiki/Tools)

BRANCH PREDICTION REALM

BRANCH PREDICTION BASICS

Why : CPUs proactively fetch instructions of potentially upcoming branches to utilize the pipeline as much as possible.

Gain if predicted correctly : If the right branch was predicted that will increase the throughput as it completed fetching a set of instructions in advance.

Penalty in case of misprediction : If the prediction was wrong, that prefetch will be a waste and the cost will be flushing the pipeline.

What are branch instructions? : Unconditional ones (jmp), conditional ones (jg, jne), call/jmp.

How : There are auxiliary hardware buffers.

Branch target buffer stores target addresses (instruction pointers) of branches. AMD uses multiple level of BTBs : L1 BTB, L2 BTB etc.

Pattern history tables track the history of results (whether it was taken or not) per branch.

A hypothetical pattern history table : T: taken, NT: not taken

CONDITIONAL MOVE INSTRUCTION

CMOV (Conditional move) instruction also computes the conditions for some additional time. Therefore they don't introduce extra load to the branch prediction mechanism. They can be used to eliminate branches.

Reference : [Intel Optimization Manual 3.4.1.1](https://en.wikichip.org/wiki/Tools)

BP METHODS : AMD PERCEPTIONS

A **perception** is basically the simplest form of machine learning. It can be considered as a linear array of weights.

Agner Fog mentions that they are good at predicting very long branches compared to 2-level adaptive branch prediction in his [microarchitecture book](https://en.wikichip.org/wiki/Tools) 3.12.

For details of perception based branch prediction : [Dynamic Branch Prediction with Perceptions by Daniel Ammerer and Calvin Lu](https://en.wikichip.org/wiki/Tools)

Note that LSD is disabled on Skylake Server CPUs. Reference : <https://en.wikichip.org/wiki/Tools>

DISABLING SPECULATIVE EXECUTION PATCHES

You can consider disabling system patches for speculative execution related vulnerabilities such as Meltdown and Spectre for performance. If that is doable in your system.

Kernel.org document : <https://en.wikichip.org/wiki/Tools>

Red Hat Enterprise documentation : <https://en.wikichip.org/wiki/Tools>

Meltdown paper : <https://en.wikichip.org/wiki/Tools>

Spectre paper : <https://en.wikichip.org/wiki/Tools>

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.

OS support

You should check your OS and CPU in combination to find out the supported states.

Linux implementation refers to them as huge pages and Windows calls them as large pages.