


MICROARCHITECTURE CHEAT SHEET

X86 CPUs & Performance



LAST UPDATE DATE : 12 OCT 2024

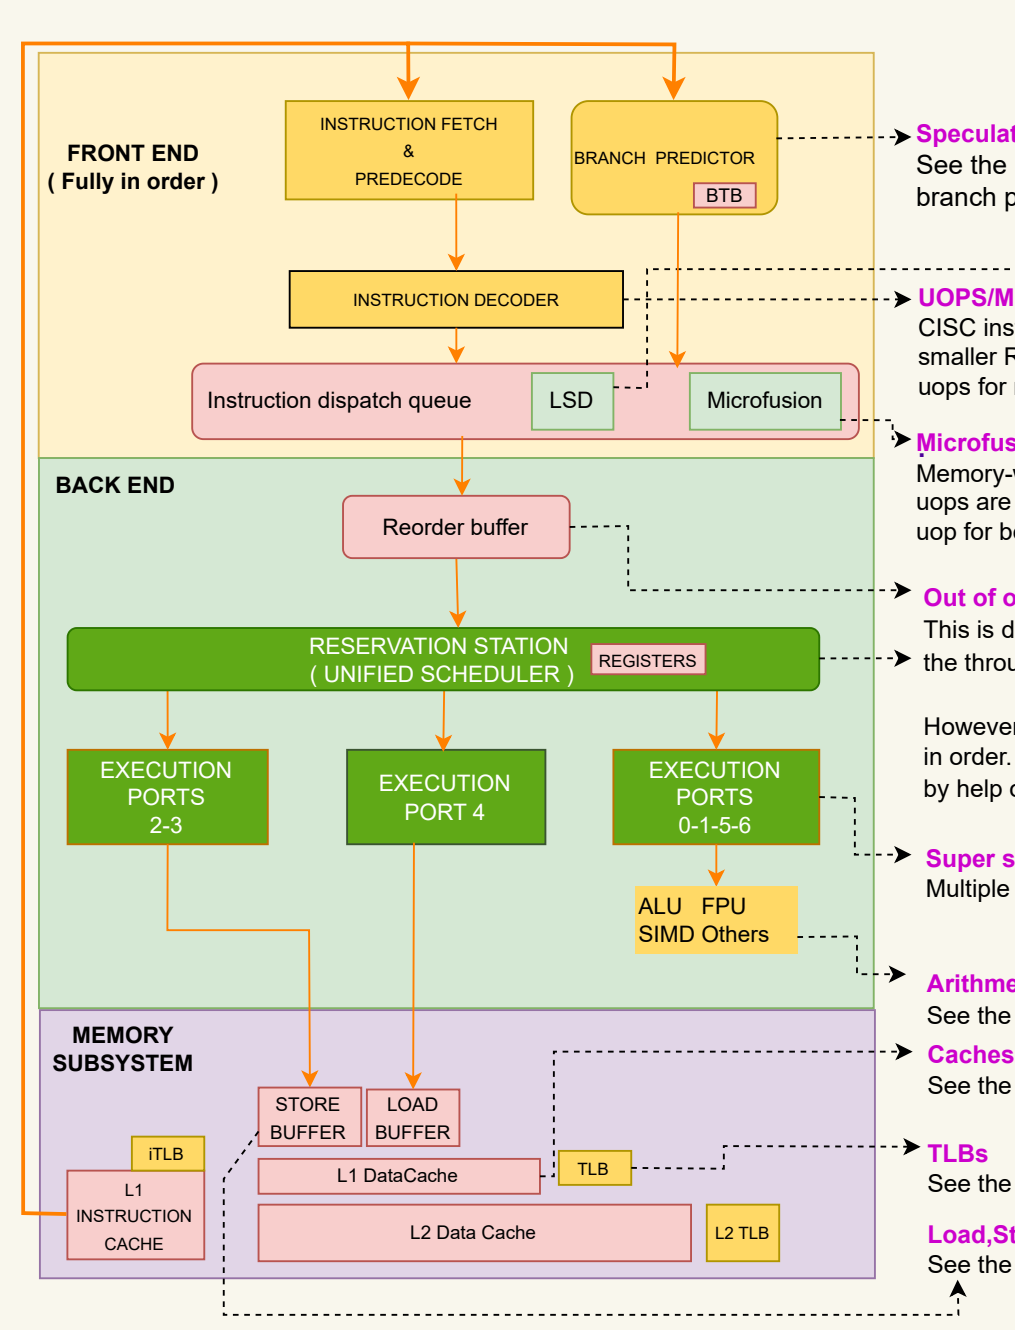
FOR LATEST VERSION : [www.github.com/akinimicroarchitecture-cheatsheet](https://github.com/akinimicroarchitecture-cheatsheet)

AUTHOR: AKIN OCAL akin_ocal@hotmail.com

PIPELINE REALM :

INSIDE AN INDIVIDUAL CORE

A SIMPLIFIED OVERVIEW (BASED ON INTEL SKYLAKE)



Based on: [Skylake server on wikipedia.org](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org)

AMD pipelines : The main difference in AMD architectures is that there are parallel pipelines for integers and floating points : [AMD Zen2 pipeline diagram on wikipedia.org](https://www.amd.com/en/technologies/amd-zen2-architecture)

PIPELINE PARALLELISM & PERFORMANCE

Pipeline Diagrams : The diagrams below in the following topics are outputs from an online microarchitecture analysis tool [LUCIA](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org) and they represent parallel execution through cycles.

Rows are multiple instructions being executed at the same time.

Columns display how instruction state changes through cycles.

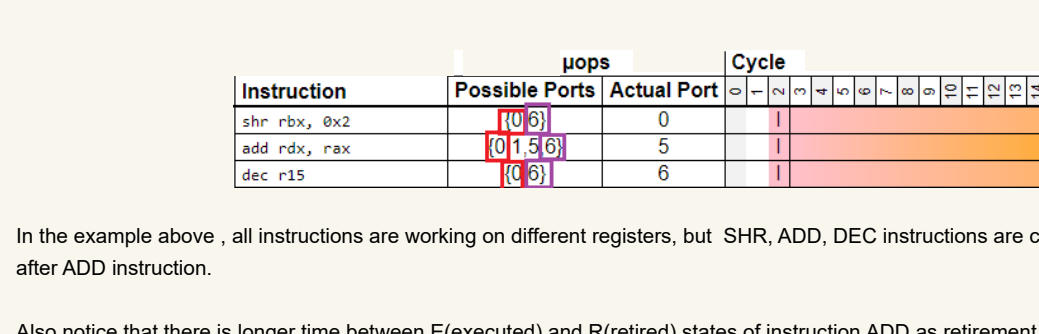
IPC : As for pipeline performance, typically IPC is used. It stands for "Instructions per cycle".

A higher IPC value usually means a better throughput.

You can measure IPC with perf : [https://perf.wiki.kernel.org/doc/2018/08/01/perf.html](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org)

Rate of retired instructions : Apart from IPC, number of retired instructions should be checked. Retired instructions are not committed/finalised as they were wrongly speculated. On the other hand executed instructions are the ones which were finalised. Therefore a high rate of retired instructions indicates low branch prediction rate.


CONTENTION FOR EXECUTION PORTS IN THE PIPELINE



In the example above, all instructions are working on different registers, but SHR, ADD, DEC instructions are competing for ports 6 and 8. SHR and DEC are getting executed after ADD instruction.

Also notice that there is longer time between (Executed) and (Retired) states of instruction ADD as retirement has to be done in-order whereas execution is out-of-order.

INSTRUCTION STALLS DUE TO DATA DEPENDENCY



In the example above, there are 2 dependency chains, each marked with a different colour. In the first red coloured one, 2 instructions are competing for the RAX register and notice that the second instruction gets executed after the first one. And the same applies to the 2nd purple pair.

Reference : [Dima Bakhtov's article](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org)

RTDSCP INSTRUCTION FOR MEASUREMENTS

TSC (time stamp counter) is a special register that counts CPU cycles. RTDSCP can be used to read the TSC value which can then be used for measurements. It can also avoid out-of-order execution effects to a degree.

It does wait until all previous instructions have executed and all previous reads are globally visible. (From [Intel's How to benchmark code execution times](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org) whitepaper has details of using RTDSCP instruction)

AMD [Zen2](https://www.amd.com/en/technologies/amd-zen2-architecture) processors : RTDSCP forces all other instructions to retire before reading the time-stamp counter

ESTIMATING INSTRUCTION LATENCIES

You can use Agner Fog's [instruction-latencies](https://www.agner-fog.com/instruction-latencies) to find out instructions' reciprocal throughput (clock cycle per instruction). As an example, reciprocal throughput of instruction RTDSCP is 32 on Skylake microarchitecture.

→ 1 cycle @ 3.5 GHz (highest frequency on Skylake) is 0.22 nanoseconds

→ 32 * 0.22 = 7.04 nanoseconds

So its resolution estimation is about 7 nanoseconds on a 3.5 GHz Skylake CPU. You have to recalculate it for different microarchitectures and clock speeds.

HYPERTHREADING / SIMULTANEOUS MULTITHREADING

Hyperthreading name is used by Intel and it is called as "Simultaneous multithreading" by AMD. In both resources including caches and execution units are shared.

Reference : Agner Fog's [processor-features](https://www.agner-fog.com/instruction-latencies) book has "multithreading" sections for each of Intel and AMD microarchitectures

Regarding using it, if your app is data-intensive, halted caches won't help. Therefore it can be disabled if via BIOS settings in general, it moves the control of resources from software to hardware and that is usually not desired for performance critical applications.

DYNAMIC FREQUENCIES

Modern CPUs employ dynamic frequency scaling which means there is a min and a max frequency per CPU core.

APCI : ACPI defines multiple power states and modern CPUs implement them. P-States are for performance and C-States are for energy efficiency.

Intel has various [hierarchy](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org) of states, and the most well known is TurboBoost. On AMD side there is [TurboCore](https://www.amd.com/en/technologies/amd-zen2-architecture). You can use those to maintain the CPU usage.

Number of active cores & SIMD AVX2512 on Intel CPUs : Intel's power management policies are complex. See the arithmetic and the multicore results as number of active cores and some of AVX2512 extensions also may affect the frequency within TurboBoost.

Varying max clock speeds on AMD CPUs : Some AMD CPUs' cores have slightly varying max frequencies. Therefore AMD CPUs have "preferred core" concept. Reference : [AMD's GDC42 presentation page13](https://www.amd.com/en/technologies/amd-zen2-architecture)

LOAD STORE REALM

LOAD AND STORE BUFFERS

Load and store buffers allow CPU to do out-of-order execution on loads and stores by decoupling speculative execution and committing the results to the cache memory.

Reference : https://en.wikipedia.org/wiki/Memory_ordering

STORE-TO-LOAD FORWARDING

Using buffers for stores and loads to support out of order execution leads to a data synchronization issue. That issue is described in en.wikipedia.org/wiki/Memory_ordering

As a solution, CPU can forward a memory store operation to a single up for better throughput.

However the results have to be output in order. So reordering at the end must be by help of the reorder buffer.

Super scalar execution : Multiple usages executed in parallel.

Arithmetic operations : See the arithmetic realm.

Caches : See the cache memory realm.

TLBs : See the virtual memory realm.

Load-store buffers : See the load-store realm.

STORE-TO-LOAD FORWARDING & LHS & PERFORMANCE

Based on [Intel Optimization Manual](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org) 3.6.4, store-to-load forwarding may improve combined latency of those 2 operations. The reason is not specified however it is potentially LHS (Load-Hit-Store) problem in which the penalty is a round trip to the cache memory :

Reference : <https://en.wikipedia.org/wiki/Load-Hit-Store>

There are several conditions for the forwarding to happen. In case of a successful forwarding, the steps 2 and 3 (a round trip to the cache) will be bypassed.

However the results have to be output in order. So reordering at the end must be by help of the reorder buffer.

The conditions for a successful forwarding and latency penalties in case of no-forwarding can be found in Agner Fog's [processor-features](https://www.agner-fog.com/instruction-latencies) book

What would happen without forwarding ? : In the past, game consoles PlayStation3 and Xbox360 had PowerPC based processors which used [out-of-order](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org) execution rather than out-of-order execution. Therefore developers had to separately handle LHS by using [out-of-order](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org) keyword and other methods : [Evan Rubin's article](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org)

ARITHMETIC REALM

ARITHMETIC INSTRUCTION LATENCIES

You can see a set of arithmetic operations table to fast to slow below.

The clock cycles are based on Agner Fog's [instruction-latencies](https://www.agner-fog.com/instruction-latencies) and Skylake architecture on 64-bit registers.

Bitwise operations, integer addition : 0.25 to 1 clock cycle

Integer multiplication : 1 to 2 clock cycles

Integer division : 25 to 30 clock cycles

FLOATING POINTS

X86 uses IEEE 754 standard for floating points. A 32-bit floating point consists of 3 parts in the memory layout. Below you can see all bits of 32-bit IEEE 754 FP number. Used <https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org>

A floating point value is calculated as : $mantissa \times 2^{exponent}$

IEEE754 also defines **denormal numbers**. They are very small / near zero numbers.

As floating points are approximations, denormal numbers are needed to avoid an undefined case of $a \times b$ but $a \neq 0$ and $b = 0$.

Without denormal the code to the right would involve a divide-by-zero exception.

Reference : [Evan Rubin's article](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org)

Based on Agner Fog's [processor-features](https://www.agner-fog.com/instruction-latencies) book, Intel CPUs have a penalty for denormal numbers. For ex. 128 clock cycles on Skylake. They also can be turned off on Intel CPUs.

As for AMD side, the recent Zen architecture CPUs seemingly don't have the same performance degradation.

X86 EXTENSIONS

X86 extensions are specialised instructions. They have various categories from [crypto](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org) to [network](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org) operations.

Intel [Instruction Guide](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org) is a good page to explore those extensions.

SSE (Streaming SIMD Extensions) is one of the most important ones that provides non-parallelism. **SMX** stands for "Single instruction multiple data". SIMD instructions use vector registers to execute more work in a single go.

In the example above, an array of 4 integers (1 to 4) are added to another array of integers (1 to 4). The result is an array of sums (1 to 4). In this example, 4 add operations are executed by a single instruction.

They play key role in complex "vectorisation" operations : [GCC auto vectorisation](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org)

Apart from arithmetic operations, they can be utilised for string operations as well : [Daniel Lemire's SIMD based ZKSN paper](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org) : <https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org>

X86 EXTENSIONS : SIMD DETAILS

The most recent SIMD instruction sets for Intel CPUs are :

AVX : Up to 256 bits

AVX2 : Up to 256 bits

AVX512 : Up to 512 bits

Recent AMD CPUs support AVX & AVX2. Only the latest Zen4 architecture supports AVX512.

As for programming, there are also wider data types. The data type diagrams below are for 32-bit operations.

...m128, 4 x 32-bit floating points

...m128, 2 x 64-bit floating points

...m128, 4 x 32-bit ints

...m128, 2 x 64-bit long ints

Note that as "SIMD instructions require more help, therefore usage of some AVX512 extensions may introduce underflowing. They should be benchmarked. For details : [Daniel Lemire's article](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org)

BRANCH PREDICTION REALM

BRANCH PREDICTION BASICS

Why : CPUs proactively fetch instructions of potentially upcoming branches to utilise the pipeline as much as possible.

Gain if predicted correctly : If the right branch was predicted that will increase the throughput as it completed fetching a set of instructions in advance.

Penalty in case of misprediction : If the prediction was wrong, that prefetch will be a waste and the cost will be flushing the pipeline.

What are branch instructions ? : Unconditional ones (jmp), conditional ones (eg. jne), callret

How : There are auxiliary hardware buffers.

Branch target buffer stores target addresses (instruction pointers) of branches. AMD uses multiple level of BTBs : L1 BTB, L2 BTB etc.

Pattern history tables track the history of results (whether it was taken or not) per branch.

CONDITIONAL MOVE INSTRUCTIONS

Conditional move instructions (for ex. CMOV) compute the conditions some additional time. However they don't introduce extra load to the branch prediction mechanism. They can be used to eliminate branches.

Reference : [Agner Fog's processor-features book 3.1.1](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org)

BP METHODS : 2-LEVEL ADAPTIVE BRANCH PREDICTION

Saturating counter as a building block

A 2-bit saturating counter can store 4 strength states.

Whenever a branch is taken it goes stronger.

And whenever a branch is not taken it goes weaker.

2 level adaptive predictor

In this method, the pattern history table keeps 2nd row and each row will have a saturating counter.

A branch history register which has the history of last n occurrences, will be used to choose which row will be used from the pattern history table.

Reference : [Agner Fog's processor-features book 3.1](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org)

BP METHODS : AMD PERCEPTIONS

A [perception](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org) is basically the simplest form of machine learning. It can be considered as a linear array of weights.

Agner Fog mentions that they are good at predicting very long branches compared to 2-level adaptive branch prediction in the [processor-features](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org) book 3.1.2.

For details of perception based branch prediction : [Dmitry Branch Prediction with Perceptions by Daniel Shewar and Calvin Lin](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org)

Intel L3D will detect a loop and stop fetching instructions to improve the frontend bandwidth.

Several conditions mentioned in [Intel Optimization Manual](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org) 3.4.2.4 :

• Loop body size up to 60 ops, with up to 15 taken branches, and up to 15 64-byte fetch lines.

• No CALL or RET.

• No mismatched stack operations (eg. more PUSH than POP).

• More than ~20 iterations.

Note that L3D is disabled on Skylake Server CPUs. Reference : <https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org>

DISABLING SPECULATIVE EXECUTION PATCHES

You can consider disabling system patches for speculative execution related vulnerabilities such as Meltdown and Spectre for performance. If that is disabled in your system. Those patches are not only [out-of-order](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org) updates but they also need OS support.

Kernel.org documentation : <https://www.kernel.org/doc/Documentation/bugzilla-guidelines.txt>

Red Hat Enterprise documentation : <https://access.redhat.com/articles/3311301>

Meltdown paper : <https://meltdownattack.com/meltdown.pdf>

Spectre paper : <https://spectreattack.com/spectre.pdf>

ESTIMATED LIMITS : HOW MANY IFs ARE TOO MANY ?

As for number of entries in BTBs, there are estimations made by stress testing the BTB with sequences of branch instructions.

Reference : [Intel Xeon Gold 6342 - roughly 4K AMD EPYC 7713 - roughly 3K](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org)

Reference : [Mark Mokroski's article on Cloudflare blog](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org)

CACHE MEMORY REALM

CACHE MEMORY VS SYSTEM MEMORY

System memory is made of DRAM cells. Cache memory on the other hand are made of SRAM cells which is much faster than DRAMs. But also they are more expensive.

DRAM used in system memories

Access time : 50-150 nanoseconds due to capacitor charge/discharge times and other steps

Cost : Cheaper in the price as it has less components

Reference : Ulrich Drepper's [What every programmer should know about memory](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org)

CACHE ORGANISATION

Caches are organised in multiple levels. As you go uppper in that hierarchy, the capacity increases. Therefore **L3** term used to indicate the last level of cache.

3 level data caches are currently the most common ones. Intel [Broadwell architecture](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org) had 4 level caches in the past. Also upcoming AMD CPUs [Zen4](https://www.amd.com/en/technologies/amd-zen2-architecture) comes with a 4 level of caches.

Cache line size : is the unit of data transfer between the cache and the system memory. It is typically 64 bytes. And the caches are organised according to the cache line size.

There is also **instruction cache** (Cache) which stores program instructions rather than data to improve throughput of CPU frontend.

In case of a cache hit, the memory is typically single digit nanoseconds. And in case of a cache miss, we need a round trip to the system memory and the latency becomes 3 digits nanoseconds.

HARDWARE AND SOFTWARE PREFETCHING

Hardware prefetchers detect patterns like **strides** (Ex. accessing to contiguous array members) and **strides** (Ex. accessing specific members in arrays of structs) and prefetch data and instruct to cache line automatically.

Developers can also use instruction **non_prelatch** to prefetch data explicitly for cases when hardware can't predict. That is called as software prefetching.

Reference for image : It is taken from [AMD's GDC42 presentation page42](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org)

N-WAY SET ASSOCIATIVITY

Why : Cache capacities are much smaller than the system memory. Moreover, software can use various regions of that address space. So if there was one to one mapping of a fully sequential memory that would lead to cache misses most of the time. Therefore there is a need for efficient mapping between the cache memory and the system memory.

How : In N-Way set associativity, caches are divided into groups of sets. And each set will have N cache blocks. The mapping information is stored in bits of addresses which has 3 parts :

used as a unique identifier per cache block

used to determine the set in a cache

used to determine the actual bytes in the target cache block

The pseudocode below shows steps for searching a single byte in the cache memory :

Get tag, set and offset from the address

For each cache in the current set (which we have just found out) :

If tag of the current block equals to tag (which we just have found out) :

need and return data using offset. It is a cache hit

If there was no matching tag, it is a cache miss

The level of associativity (the number of ways) is a trade off between the search time and the amount of system memory we can map.

BYPASSING THE CACHE : NON-TEMPORAL STORES & WRITE-COMBINING

Temporal data is data that will be accessed in a short period of time. The term **non-temporal** data indicates that data will not be accessed any soon (cold data). If the amount of non-temporal data is excessive in the cache, that is called as **cache pollution**. Non-temporal store instructions are introduced for this problem and they store data directly to the system memory by bypassing the cache.

Write combining buffers are used with non-temporal stores. CPU will try to fill a whole cache line (typically 64bytes) before committing to the system memory and only will send to the system memory when that buffer is filled. That is for reducing the load on the bus between the CPU and the system memory.

DMA mechanism doesn't require CPU involvement. Through mechanism initiated by CPU, therefore CPU support needed.

DCA bypasses the system memory and can transfer to directly LLC of CPUs that support the technology.

Intel refers to this technology as DQO (Direct I/O).

Reference : [Intel documentation](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org)

DIRECT CACHE ACCESS

Modern NICs come with a DMA (Direct Memory Access) engine and can transfer data directly to driver ring buffers which reside on the system memory.

DMA mechanism doesn't require CPU involvement. Through mechanism initiated by CPU, therefore CPU support needed.

DCA bypasses the system memory and can transfer to directly LLC of CPUs that support the technology.

Intel refers to this technology as DQO (Direct I/O).

Reference : [Intel documentation](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org)

VIRTUAL MEMORY REALM

VIRTUAL MEMORY ORGANISATION

Why virtual memory ?

Because cumulative memory requirement of multiple processes in an OS can be more than the system capacity.

It is basically for sharing memory resources between multiple processes.

It also provides security by isolating processes.

Pages : Minimum addressable virtual space that can be requested from OS.

Typically 4KB. But there may be extra allocations on top of page size. For ex. Windows has 4KB page allocation granularity : [Raymond Chen's article](https://www.intel.com/content/dam/develop/external/us/en/documents/skylake-server-on-wikipedia.org)

Swapping : In case of out of memory, process memory will be evicted to the disc.

Page faults : Happens when the page is not on physical memory but on the swap file which is on the hard drive.

VIRTUAL MEMORY ADDRESS TRANSLATION

Address translation & Page table